



Introduction to Numerical Computing with Numpy

Use Restrictions:

Use only permitted under license or agreement. Copying, sharing, redistributing or other unauthorized use is strictly prohibited. The Enthought Training Materials (ETM) are provided for the individual and sole use of the paid attendee of the class ("Student") for which the Training Services are provided. The virtual training sessions may not be shared, reproduced, transmitted, or retransmitted in any form or by any means, electronic or mechanical, including photocopying, recording, or any information storage and retrieval system. Furthermore, neither Customer nor any Student shall:

Copy, disclose, transfer or distribute ETM to any party in any form. Remove, modify or obscure any copyright, trademark, legal notices or other proprietary notations in ETM. Make derivative works of ETM or combine ETM or any part of ETM with any other works. Use ETM in any manner that could be detrimental to Enthought. • 2001-2022, Enthought, Inc. All Rights Reserved. All trademarks and registered trademarks are the property of their respective owners.

Enthought, Inc. 200 W Cesar Chavez Suite 202 Austin, TX 78701
www.enthought.com

Q3-2022
letter
3.5.5



Introduction to Numerical Computing with Numpy

Enthought, Inc.
www.enthought.com

Introduction 1

NumPy 2

- Introducing NumPy Arrays 4

- Multi-Dimensional Arrays 8

- Slicing/Indexing Arrays 10

- Fancy Indexing 16

- Creating arrays 19

- Array Creation Functions 21

- Array Calculation Methods 24

- Array Broadcasting 33

- Universal Function Methods 41

- The array data structure 48

- Creating arrays 51

Closing words 56

Appendix 57

About Enthought 58

Introduction to Numerical Computing with NumPy

SciPy Conference Tutorial
2022

1


NumPy

The Standard Numerical Library
for Python

2

Syllabus

- 1. Introducing Arrays
- 2. Indexing and Slicing
- 3. Creating Arrays
- 4. Array Calculations
- 5. The Array Data Structure
- 6. Structure Operations



NumPy

Introducing Arrays

Introducing NumPy Arrays

SIMPLE ARRAY CREATION

```
>>> a = np.array([0, 1, 2, 3])
>>> a
array([0, 1, 2, 3])
```

CHECKING THE TYPE

```
>>> type(a)
numpy.ndarray
```

NUMERIC "TYPE" OF ELEMENTS

```
>>> a.dtype
dtype('int32')
```

NUMBER OF DIMENSIONS

```
>>> a.ndim
1
```

ARRAY SHAPE

```
# Shape returns a tuple listing the
# length of the array along each
# dimension.
>>> a.shape
(4,)
```

BYTES PER ELEMENT

```
>>> a.itemsize
4
```

BYTES OF MEMORY USED

```
# Return the number of bytes used by
# the data portion of the array.
>>> a.nbytes
16
```

5

Array Operations

SIMPLE ARRAY MATH

```
>>> a = np.array([1, 2, 3, 4])
>>> b = np.array([2, 3, 4, 5])
>>> a + b
array([3, 5, 7, 9])
```

```
>>> a * b
array([ 2,  6, 12, 20])
```

```
>>> a ** b
array([ 1,  8, 81, 1024])
```

MATH FUNCTIONS

```
# create array from 0.0 to 10.0
>>> x = np.arange(11.0)
```

```
# multiply entire array by scalar
# value
>>> c = (2.0 * np.pi) / 10.0
>>> c
0.62831853071795862
>>> c * x
->>- E
```

```
# in-place operations
>>> x *= c
>>> x
->>- E
```

```
# apply functions to array
>>> y = np.sin(x)
```




NumPy defines these constants:
pi = 3.14159265359
e = 2.71828182846

6


Setting Array Elements

ARRAY INDEXING

```
>>> a[0]
0
```



```
>>> a[0] = 10
>>> a
array([10, 1, 2, 3])
```



BEWARE OF TYPE COERCION

```
>>> a.dtype
dtype('int32')


# assigning a float into an int32
# array truncates the decimal part
>>> a[0] = 10.6
>>> a
array([10, 1, 2, 3])

# fill has the same behavior
>>> a.fill(-4.8)
>>> a
array([-4, -4, -4, -4])
```

Multi-Dimensional Arrays


MULTI-DIMENSIONAL ARRAYS

```
>>> a = np.array([[ 0, 1, 2, 3],
...               [10, 11, 12, 13]])
>>> a
array([[ 0, 1, 2, 3],
       [10, 11, 12, 13]])
```




SHAPE (ROWS, COLUMNS)

```
>>> a.shape
(2, 4)
```




ELEMENT COUNT

```
>>> a.size
8
```



NUMBER OF DIMENSIONS


```
>>> a.ndim
2
```



ARRAY SHAPE


```
>>> a[1, 3]
13

>>> a[1, 3] = -1
>>> a
array([[ 0, 1, 2, 3],
       [10, 11, 12, -1]])
```



ADDRESS SECOND (ONETH) ROW USING SINGLE INDEX

```
>>> a[1]
array([10, 11, 12, -1])
```



Formatting Numeric Display

DEFAULT FORMATTING

```
>>> a = np.arange(1.0, 3.0, 0.5)
>>> a
array([1. , 1.5, 2. , 2.5])

>>> a * np.pi
array([3.14159265, 4.71238898, 6.28318531,
       7.85398163])

>>> a * np.pi * 1e8
array([3.14159265e+08, 4.71238898e+08,
       6.28318531e+08, 7.85398163e+08])

>>> a * np.pi * 1e-6
array([3.14159265e-06, 4.71238898e-06,
       6.28318531e-06, 7.85398163e-06])
```

USER FORMATTING

```
# set precision
>>> np.set_printoptions(
    precision=2)


>>> a
array([1. , 1.5, 2. , 2.5])

>>> a * np.pi
array([3.14, 4.71, 6.28, 7.85])

>>> a * np.pi * 1e8
array([3.14e+08, 4.71e+08, 6.28e+08,
       7.85e+08])

>>> a * np.pi * 1e-6
array([3.14e-06, 4.71e-06, 6.28e-06,
       7.85e-06])

# suppress scientific notation
>>> np.set_printoptions(
    suppress=True)
>>> a * np.pi * 1e-6
array([0. , 0. , 0. , 0.])
```



NumPy

Indexing and Slicing

Slicing

var[lower: upper: step]

- “ Extracts a portion of a sequence by specifying a lower and upper bound.
- “ The lower-bound element is included, but the upper-bound element is not included.
- “ Mathematically: [lower, upper). The step value specifies the stride between elements.

SLICING ARRAYS

```
#           -5 -4 -3 -2 -1
# indices:  0  1  2  3  4
>>> a = np.array([10, 11, 12, 13, 14])

# [10, 11, 12, 13, 14]
>>> a[1:3]
array([11, 12])

# negative indices work also
>>> a[1:-2]
array([11, 12])
>>> a[-4:3]
array([11, 12])
```

OMITTING INDICES

```
# omitted boundaries are assumed to be
# the beginning (or end) of the array

# grab first three elements
>>> a[:3]
array([10, 11, 12])

# grab last two elements
>>> a[-2:]
array([13, 14])

# every other element
>>> a[::2]
array([10, 12, 14])
```

Array Slicing

SLICING WORKS MUCH LIKE STANDARD PYTHON SLICING


```
>>> a[0, 3:5]
array([3, 4])

>>> a[4:, 4:]
array([[44, 45],
       [54, 55]])

>>> a[:, 2]
array([2, 12, 22, 32, 42, 52])
```

SLICING ARRAYS

```
>>> a[2::2, ::2]
array([[20, 22, 24],
       [40, 42, 44]])
```



Assigning to a Slice

Slices are references to locations in memory.
These memory locations can be used in assignment operations.

```
>>> a = np.array([0, 1, 2, 3, 4])  
# slicing the last two elements returns the data there  
>>> a[-2:]  
array([3, 4])  
  
# we can insert an iterable of length two  
>>> a[-2:] = [-1, -2]  
>>> a  
array([0, 1, 2, -1, -2])  
  
# or a scalar value  
>>> a[-2:] = 99  
>>> a  
array([0, 1, 2, 99, 99])
```

13

Give it a try!



Create the array below with the command

```
a = np.arange(25).reshape(5, 5)
```

and extract the slices as indicated.

	0	1	2	3	4
0	0	1	2	3	4
1	5	6	7	8	9
2	10	11	12	13	14
3	15	16	17	18	19
4	20	21	22	23	24

Sliced Arrays Share Data

Arrays created by slicing share data with the originating array.
Changing values in a slice also changes the original array.

```
>>> a = np.array([0, 1, 2, 3, 4])

# create a slice containing two elements of a
>>> b = a[2:4]
>>> b
array([2, 3])
>>> b[0] = 10

# changing b changed a!
>>> a
array([ 0,  1, 10,  3,  4])
>>> np.shares_memory(a, b)
True
```

Fancy Indexing

INDEXING BY POSITION

```
>>> a = np.arange(0, 80, 10)


# fancy indexing
>>> indices = [1, 2, -3]
>>> y = a[indices]
>>> y
array([10, 20, 50])

# this also works with setting
>>> a[indices] = 99
>>> a
array([0, 99, 99, 30, 40, 99, 60, 70])
```

INDEXING WITH BOOLEANS

```
# manual creation of masks
>>> mask = np.array(
...     [0, 1, 1, 0, 0, 1, 0, 0],
...     dtype=bool)

# fancy indexing
>>> y = a[mask]
>>> y
array([99, 99, 99])
```



Fancy Indexing in 2-D

```
>>> a[[0, 1, 2, 3, 4],  
...   [1, 2, 3, 4, 5]]  
array([ 1, 12, 23, 34, 45])
```

```
>>> a[3:, [0, 2, 5]]  
array([[30, 32, 35],  
       [40, 42, 45],  
       [50, 52, 55]])
```

```
>>> mask = np.array(  
...     [1, 0, 1, 0, 0, 1],  
...     dtype=bool)  
>>> a[mask, 2]  
array([2, 22, 52])
```

	0	1	2	3	4	5
0	0	1	2	3	4	5
1	10	11	12	13	14	15
2	20	21	22	23	24	25
3	30	31	32	33	34	35
4	40	41	42	43	44	45
5	50	51	52	53	54	55



Unlike slicing, fancy indexing creates copies instead of a view into original array.

Give it a try!



1. Create the array below with

```
a = np.arange(25).reshape(5, 5)  
and extract the elements indicated in blue.
```

2. Extract all the numbers divisible by 3 using a boolean mask.

	0	1	2	3	4
0	0	1	2	3	4
1	5	6	7	8	9
2	10	11	12	13	14
3	15	16	17	18	19
4	20	21	22	23	24

```
000
arange()
linspace()
array()
zeros()
ones()
```

NumPy

Creating Arrays

Array Constructor Examples

FLOATING POINT ARRAYS


```
# Default to double precision
>>> a = np.array([0, 1.0, 2, 3])
>>> a.dtype
dtype('float64')
>>> a.nbytes
32
```

REDUCING PRECISION

```
>>> a = np.array([0, 1., 2, 3],
...               dtype='float32')
>>> a.dtype
dtype('float32')
>>> a.nbytes
16
```

ARRAY SHAPE

```
>>> a = np.array([0, 1, 2, 3],
...               dtype='uint8')
>>> a.dtype
dtype('uint8')
>>> a.nbytes
4
```



Base 2	Base 10
00000000	-> 0 = 0*2**0 + 0*2**1 + ... + 0*2**7
00000001	-> 1 = 1*2**0 + 0*2**1 + ... + 0*2**7
00000010	-> 2 = 0*2**0 + 1*2**1 + ... + 0*2**7
...	
11111111	-> 255 = 1*2**0 + 1*2**1 + ... + 1*2**7

Array Creation Functions

ARANGE

```
arange([start, ] stop[, step],
       dtype=None)
```

Creates an array of values in the range [start,stop) with the specified step value. Allows non-integer values for start, stop, and step. Default dtype is derived from the start, stop, and step values.

```
>>> np.arange(4)
array([0, 1, 2, 3])
>>> np.arange(0, 2*pi, pi/4)
array([ 0.000, 0.785, 1.571, 2.356, 3.142,
        3.927, 4.712, 5.497])
```


```
>>> np.arange(1.5, 2.1, 0.3)
array([ 1.5,  1.8,  2.1])
```

ONES, ZEROS

```
ones(shape, dtype='float64')
zeros(shape, dtype='float64')
```

shape is a number or sequence specifying the dimensions of the array. If dtype is not specified, it defaults to float64.

```
>>> np.ones((2, 3),
...          dtype='float32')
array([[ 1.,  1.,  1.],
       [ 1.,  1.,  1.]])
>>> np.zeros(3)
array([ 0.,  0.,  0.])
```



Array Creation Functions (cont'd)

IDENTITY

```
# Generate an n by n identity array.
# The default dtype is float64.
>>> a = np.identity(4)
>>> a
array([[ 1.,  0.,  0.,  0.],
       [ 0.,  1.,  0.,  0.],
       [ 0.,  0.,  1.,  0.],
       [ 0.,  0.,  0.,  1.]])
>>> a.dtype
dtype('float64')
>>> np.identity(4, dtype=int)
array([[ 1,  0,  0,  0],
       [ 0,  1,  0,  0],
       [ 0,  0,  1,  0],
       [ 0,  0,  0,  1]])
```

EMPTY AND FULL

```
# empty(shape, dtype=float64,
#        order='C')
>>> np.empty(2)
array([ 1.78021120e-306,
        6.95357225e-308])

# array filled with 5.0
>>> a = np.full(2, 5.0)
array([ 5.,  5.])

# alternative approaches
# (slower)
>>> a = np.empty(2)
>>> a.fill(4.0)
>>> a
array([ 4.,  4.])
>>> a[:] = 3.0
>>> a
array([ 3.,  3.])
```

Array Creation Functions (cont'd)



Linspace

```
# Generate N evenly spaced elements
# between (and including) start and
# stop values.
>>> np.linspace(0, 1, 5)
array([0., 0.25, 0.5, 0.75, 1.0])
```

Logspace

```
# Generate N evenly spaced elements on
# a log scale between base**start and
# base**stop (default base=10)
>>> np.logspace(0, 1, 5)
array([1., 1.78, 3.16, 5.62, 10.])
```

Arrays from/to text files


```
BEGINNING OF THE FILE
% Day, Month, Year, Skip, Avg Power
01, 01, 2000, x876, 13 % crazy day!
C1 0; : @ 4-B1 fl- : >0
04, 01, 2000, xfed, 55
Data.txt
```

```
# loadtxt() automatically generates an
# array from the txt file
```

```
arr = np.loadtxt('Data.txt',
...             skiprows=1,
...             dtype=int, delimiter=",",
...             usecols = (0, 1, 2, 4),
...             comments = "%")
```

```
# Save an array into a txt file
np.savetxt('filename.txt', arr)
```

© 2008-2022 Enthought, Inc. 23



NumPy


Array Calculation Methods

Computations with Arrays

Rule 1:	Operations between multiple array objects are first checked for proper shape match.
Rule 2:	Mathematical operators (+ - * / exp, log, ...) apply element by element, on the values.
Rule 3:	Reduction operations (mean, std, skew, kurt, sum, prod, ...) apply to the whole array, unless an axis is specified.
Rule 4:	Missing values propagate unless explicitly ignored (nanmean, nansum, ...).

Multi-Dimensional Arrays

VISUALIZING MULTI-DIMENSIONAL ARRAYS



Array Calculation Methods


SUM METHOD

```
# Methods act on data stored in the array
>>> a = np.array([[1, 2, 3],
                  [4, 5, 6]])

# .sum() defaults to adding up all the
# values in an array.
>>> a.sum()
21

# supply the keyword axis to sum along the
# 0th axis
>>> a.sum(axis=0)
array([5, 7, 9])

# supply the keyword axis to sum along the
# last axis
>>> a.sum(axis=-1)
array([ 6, 15])
```



© 2008-2022 Enthought, Inc. 27

Other Operations on Arrays

SUM FUNCTION

```
# Functions work on data passed to it
>>> a = np.array([[1, 2, 3],
                  [4, 5, 6]])

# sum() defaults to adding up all
# values in an array.
>>> np.sum(a)
21

# supply an axis argument to sum along
# a specific axis
>>> np.sum(a, axis=0)
array([5, 7, 9])
```

OTHER METHODS AND FUNCTIONS

Mathematical functions

- `sum`, `prod`
- `min`, `max`, `argmin`, `argmax`
- `ptp` (max - min)

Statistics

- `mean`, `std`, `var`

Truth value testing

- `any`, `all`

See the NumPy appendix for more.

Min/Max

MIN	ARGMIN/MAX
<pre>>>> a = np.array([[2, 3], [0, 1]]) # Prefer NumPy functions to builtins when # working with arrays >>> np.min(a) 0 # Most NumPy reducers can be used as # methods as well as functions >>> a.min() 0</pre>	<pre># Many tasks (like optimization) are # interested in the location of a min/max, # not the value >>> a.argmax() 1 # arg methods return the location in a 1D, # flattened version of the original array >>> np.argmin(a) 2</pre>
MAX	UNRAVELING
<pre># Use the axis keyword to find max values # for one dimension >>> a.max(axis=0) array([2, 3]) # as a function >>> np.max(a, axis=1) array([3, 2])</pre>	<pre># NumPy includes a function to un-flatten # 1D locations >>> np.unravel_index(... a.argmax(), a.shape) (0, 1)</pre>

29

Where



COORDINATE LOCATIONS	CONDITIONAL ARRAY CREATION
<pre># NumPy's where function has two # distinct uses. One is to provide # coordinates from masks. >>> a = np.arange(-2, 2) ** 2 >>> a array([4, 1, 0, 1]) >>> mask = a < 1 >>> mask array([False, True, True, False]) # Coordinates are returned as a tuple # of arrays, one for each axis. >>> np.where(mask) (array([1, 2]),)</pre>	<pre># Where can also be used to construct a new # array by choosing values from other # arrays of the same shape. >>> positives = np.arange(1, 5) >>> negatives = -positives >>> np.where(mask, positives, ... negatives) array([1, -2, 3, -4]) # Or from scalar values. This can be useful # for recoding arrays. >>> np.where(mask, 1, 0) array([1, 0, 1, 0]) # Or from both. >>> np.where(mask, positives, 0) array([1, 0, 3, 0])</pre>

30

Statistics Array Methods

MEAN

```
>>> a = np.array([[1, 2, 3],
...               [4, 5, 6]])

# mean value of each column
>>> a.mean(axis=0)
array([ 2.5,  3.5,  4.5])
>>> np.mean(a, axis=0)
array([ 2.5,  3.5,  4.5])
```

STANDARD DEV./VARIANCE

```
# Standard Deviation
>>> a.std(axis=0)
array([ 1.5,  1.5,  1.5])
# For sample, set ddof=1
>>> a.std(axis=0, ddof=1)
array([ 2.12,  2.12,  2.12])

# variance
>>> a.var(axis=0)
array([2.25,  2.25,  2.25])
>>> np.var(a, axis=0)
array([2.25,  2.25,  2.25])
```

Give it a try!



Create the array below with

```
a = np.arange(-15, 15).reshape(5, 6) ** 2
```

and compute:

- 1. The maximum of each row
- 2. The mean of each column
- 3. The position of the overall minimum

	0	1	2	3	4	5
0	225	196	169	144	121	100
1	81	64	49	36	25	16
2	9	4	1	0	1	4
3	9	16	25	36	49	64
4	81	100	121	144	169	196

a	b	y
0	0	0
10	10	10
20	20	20
30	30	30

+

=

y

0

1

2

10

20

30

31

32

NumPy

Array Broadcasting

© 2008-2022 Enthought, Inc. 33

Array Broadcasting

NumPy arrays of different dimensionality can be combined in the same expression. Arrays with smaller dimension are broadcasted to match the larger arrays, *without copying data*.

Broadcasting has two rules.

RULE 1: PREPEND ONES TO SMALLER ARRAY'S SHAPE

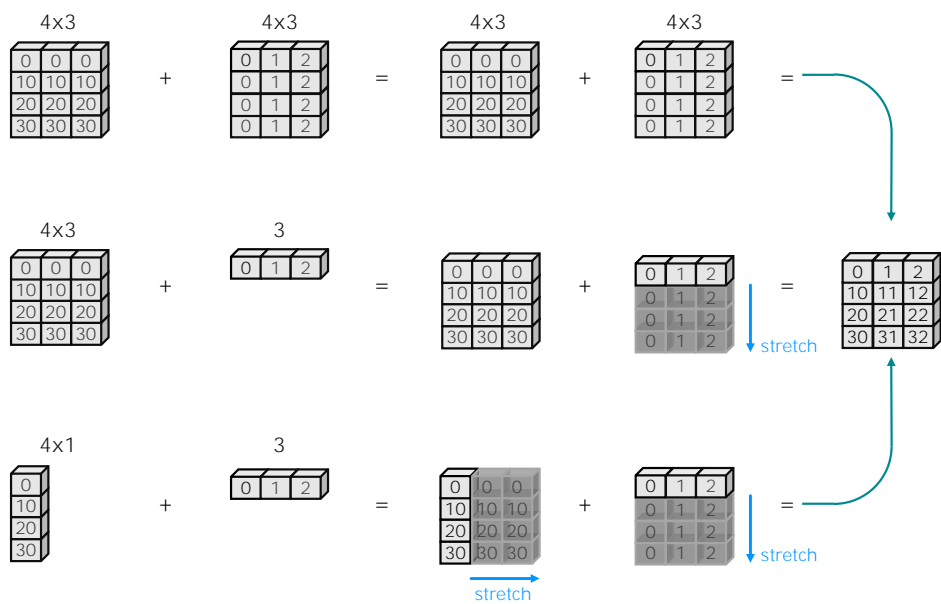
```
>>> import numpy as np
>>> a = np.ones((3, 5)) # a.shape == (3, 5)
>>> b = np.ones((5,)) # b.shape == (5,)
>>> b.reshape(1, 5) # result is a (1,5)-shaped array
>>> b[np.newaxis, :] # equivalent, more concise
```

RULE 2: DIMENSIONS OF SIZE 1 ARE REPEATED WITHOUT COPYING

```
>>> c = a + b # c.shape == (3, 5)
# is logically equivalent to...
>>> tmp_b = b.reshape(1, 5)
>>> tmp_b_repeat = tmp_b.repeat(3, axis=0)
>>> c = a + tmp_b_repeat
# But broadcasting makes no copies of "b"s data!
```

© 2008-2022 Enthought, Inc. 34

Array Broadcasting




Broadcasting Rules

The trailing axes of either arrays must be 1 or both must have the same size for broadcasting to occur.

Otherwise, a


"ValueError: shape mismatch: objects cannot be broadcast to a single shape" exception is thrown.





Broadcasting in Action

```
>>> a = array([0, 10, 20, 30])
>>> b = array([0, 1, 2])
>>> y = a[:, newaxis] + b
```



© 2008-2022 Enthought, Inc. 37

Application: Distance from Center

```
ooo
>>> import matplotlib.pyplot as plt
>>> a = np.linspace(0, 1, 15) - 0.5
>>> b = a[:, np.newaxis] # b.shape == (15, 1)
>>> dist2 = a**2 + b**2 # broadcasting sum.
>>> dist = np.sqrt(dist2)
>>> plt.imshow(dist)
>>> plt.colorbar()
```

© 2008-2022 Enthought, Inc. 38

Broadcasting's Usefulness

Broadcasting can often be used to replace needless data replication inside a NumPy array expression.

`np.meshgrid()` use newaxis appropriately in broadcasting expressions.

`np.repeat()` broadcasting makes repeating an array along a dimension of size 1 unnecessary.

MESHGRID: COPIES DATA

```
>>> x, y = \
...     np.meshgrid([1, 2],
...                  [3, 4, 5])
>>> z = x + y
```

BROADCASTING: NO COPIES


```
>>> x = np.array([1, 2])
>>> y = np.array([3, 4, 5])
>>> z = x[np.newaxis, :] \
...     + y[:, np.newaxis]
```

Broadcasting Indices


Broadcasting can also be used to get a slice of a 3-D (or any other shape) array. This is a very powerful feature of indexing.


```
>>> data_cube = np.arange(27).reshape(3, 3, 3)
>>> yi, xi = np.meshgrid(np.arange(3), np.arange(3),
...                      sparse=True)
>>> zi = np.array([[0, 1, 2],
...                [1, 1, 2],
...                [2, 2, 2]])
>>> horizon = data_cube[xi, yi, zi]
```

Indices



Selected Data





NumPy

Universal Function Methods

© 2008-2022 Enthought, Inc. 41

Universal Function Methods

The mathematical, comparative, logical, and bitwise operators *op* that take two arguments (binary operators) have special methods that operate on arrays:

```
>>> op.reduce(a, axis=0)
>>> op.accumulate(a, axis=0)
>>> op.outer(a, b)
>>> op.reduceat(a, indices)
```

© 2008-2022 Enthought, Inc. 42

op. reduce()

op. reduce(a) applies **op** to all the elements in a 1-D array `a` reducing it to a single value.

For example:

```
y = add.reduce(a)
```

ADD EXAMPLE

```
>>> a = np.array([1, 2, 3, 4])
>>> np.add.reduce(a)
10
```

STRING LIST EXAMPLE

```
>>> a = np.array(['ab', 'cd', 'ef'],
...               dtype='object')
>>> np.add.reduce(a)
'abcdef'
```

LOGICAL OP EXAMPLES


```
>>> a = np.array([1, 1, 0, 1])
>>> np.logical_and.reduce(a)
False
>>> np.logical_or.reduce(a)
True
```

op. reduce()

For multidimensional arrays, `op. reduce(a, axis)` applies `op` to the elements of `a` along the specified `axis`. The resulting array has dimensionality one less than `a`. The default value for `axis` is 0.


SUMMING UP EACH ROW

```
>>> a = np.arange(3) + np.arange(0, 40,
...               10).reshape(-1, 1)
>>> np.add.reduce(a, 1)
array([ 3, 33, 63, 93])
```



SUM COLUMNS BY DEFAULT

```
>>> np.add.reduce(a)
array([60, 64, 68])
```



op. accumulate()

op. accumulate(a) creates a new array containing the intermediate results of the reduce operation at each element in a.

For example:

```
y = add.accumulate(a)
```

ADD EXAMPLE

```
>>> a = np.array([1, 2, 3, 4])
>>> np.add.accumulate(a)
array([ 1,  3,  6, 10])
```

STRING LIST EXAMPLE

```
>>> a = np.array(['ab', 'cd', 'ef'],
...               dtype='object')
>>> np.add.accumulate(a)
array(['ab', 'abcd', 'abcdef'],
      dtype=object)
```

LOGICAL OP EXAMPLES

```
>>> a = np.array([1, 1, 0])
>>> np.logical_and.accumulate(a)
array([True,  True, False])
>>> np.logical_or.accumulate(a)
array([True,  True,  True])
```

op. reduceat()


op. reduceat(a, indices) applies op to ranges in the 1-D array a defined by the values in indices. The resulting array has the same length as indices.

For example:

```
y = add.reduceat(a, indices)
```

EXAMPLE


```
>>> a = np.array([0, 10, 20, 30, 40, 50])
>>> indices = np.array([1, 4])
>>> np.add.reduceat(a, indices)
array([60, 90])
```



For multidimensional arrays, reduceat() is always applied along the last axis (sum of rows for 2-D arrays). This is different from the default for reduce() and accumulate().

op. outer()

op. outer(a, b) forms all possible combinations of elements between a and b using op. The shape of the resulting array results from concatenating the shapes of a and b. (Order matters.)




```
>>> np.add.outer(a, b)
```

a[0]+b[0]	a[0]+b[1]	a[0]+b[2]
a[1]+b[0]	a[1]+b[1]	a[1]+b[2]
a[2]+b[0]	a[2]+b[1]	a[2]+b[2]
a[3]+b[0]	a[3]+b[1]	a[3]+b[2]

```
>>> np.add.outer(b, a)
```


b[0]+a[0]	b[0]+a[1]	b[0]+a[2]	b[0]+a[3]
b[1]+a[0]	b[1]+a[1]	b[1]+a[2]	b[1]+a[3]
b[2]+a[0]	b[2]+a[1]	b[2]+a[2]	b[2]+a[3]

Memory Block

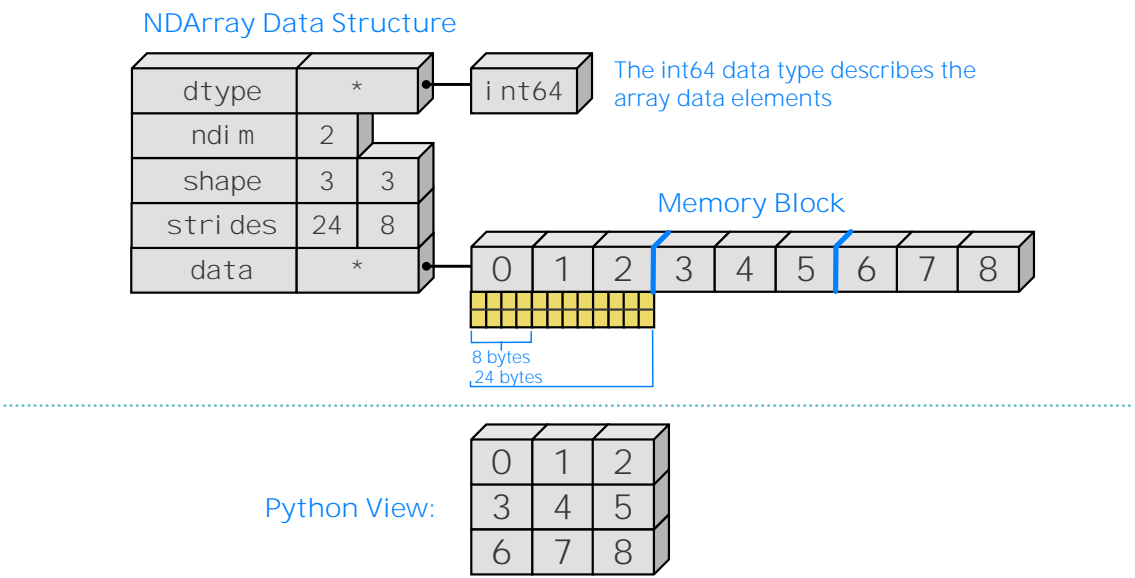



NumPy
The Array Data Structure

Array Data Structure



Array Data Structure






NumPy

Structure Operations

Operations on the Array Structure

Operations that only affect the array structure, not the data, can usually be executed without copying memory.




This **is not** a new copy of the data.
The original data **does not** get reordered.

Transpose

TRANPOSE


```
>>> a = np.array([[0, 1, 2],
...               [3, 4, 5]])
>>> a.shape
(2, 3)

# Transpose swaps the order of axes.
>>> a.T
array([[0, 3],
       [1, 4],
       [2, 5]])
>>> a.T.shape
(3, 2)
```



TRANPOSE RETURNS VIEWS

```
# Transpose does not move values around in
# memory. It only changes the order of
# "strides" in the array
>>> a.strides
(24, 8)
>>> a.T.strides
(8, 24)
```



Reshaping Arrays

RESHAPE


```
>>> a = np.array([[0, 1, 2],
...               [3, 4, 5]])
...
# Return a new array with a different shape
# (a view where possible)
>>> a.reshape(3, 2)
array([[0, 1],
       [2, 3],
       [4, 5]])

# Reshape cannot change the number of
# elements in an array
>>> a.reshape(4, 2)
ValueError: total size of new array must be
unchanged
```

SHAPE

```
>>> a = np.arange(6)
>>> a
array([0, 1, 2, 3, 4, 5])
>>> a.shape
(6,)

# Reshape array in-place to 2x3
>>> a.shape = (2, 3)
>>> a
array([[0, 1, 2],
       [3, 4, 5]])
```



Flattening Arrays


FLATTEN (SAFE)

a. `flatten()` converts a multi-dimensional array into a 1-D array. The new array is a copy of the original data.

```
# Create a 2D array
>>> a = np.array([[0, 1],
...               [2, 3]])

# Flatten out elements to 1D
>>> b = a.flatten()
>>> b
array([0, 1, 2, 3])

# Changing b does not change a
>>> b[0] = 10
>>> b
array([10, 1, 2, 3])
>>> a
array([[0, 1],
       [2, 3]])
```




RAVEL (EFFICIENT)

a. `ravel()` is the same as `flatten()`, but returns a reference (or view) of the array if possible (i.e., the memory is contiguous). Otherwise the new array copies the data. `np.ravel()` can be applied to non-array objects.

```
# Flatten out elements to 1-D
>>> b = a.ravel()
>>> b
array([0, 1, 2, 3])

# Changing b does change a
>>> b[0] = 10
>>> b
array([10, 1, 2, 3])
>>> a
array([[10, 1],
       [ 2, 3]])
```



Stay in touch!



Enthought

Enthought
Media

@enthought




EuroSciPy

Please complete the online survey!
(link on course web page)

Appendix

Additional Material

We Partner with Industry Leaders



ABOUT ENTHOUGHT

Foundation for Success

Global, privately-owned company with headquarters in Austin

Offices in Houston, Texas; Cambridge, United Kingdom; Zurich, Switzerland; and Tokyo, Japan

Deep roots in the Python community of scientists and engineers

80% of employees have graduate degrees or PhDs in science and engineering

Partner with Enthought

Science is at the core of your business. And ours.

Designed by scientists, for scientists

The Enthought Approach is proven.

Catalyzing innovation by transforming technology, people, and processes

• • • • •

Deep scientific expertise and computational knowledge

We empower scientists.

Equipping scientists with digital skills and technology